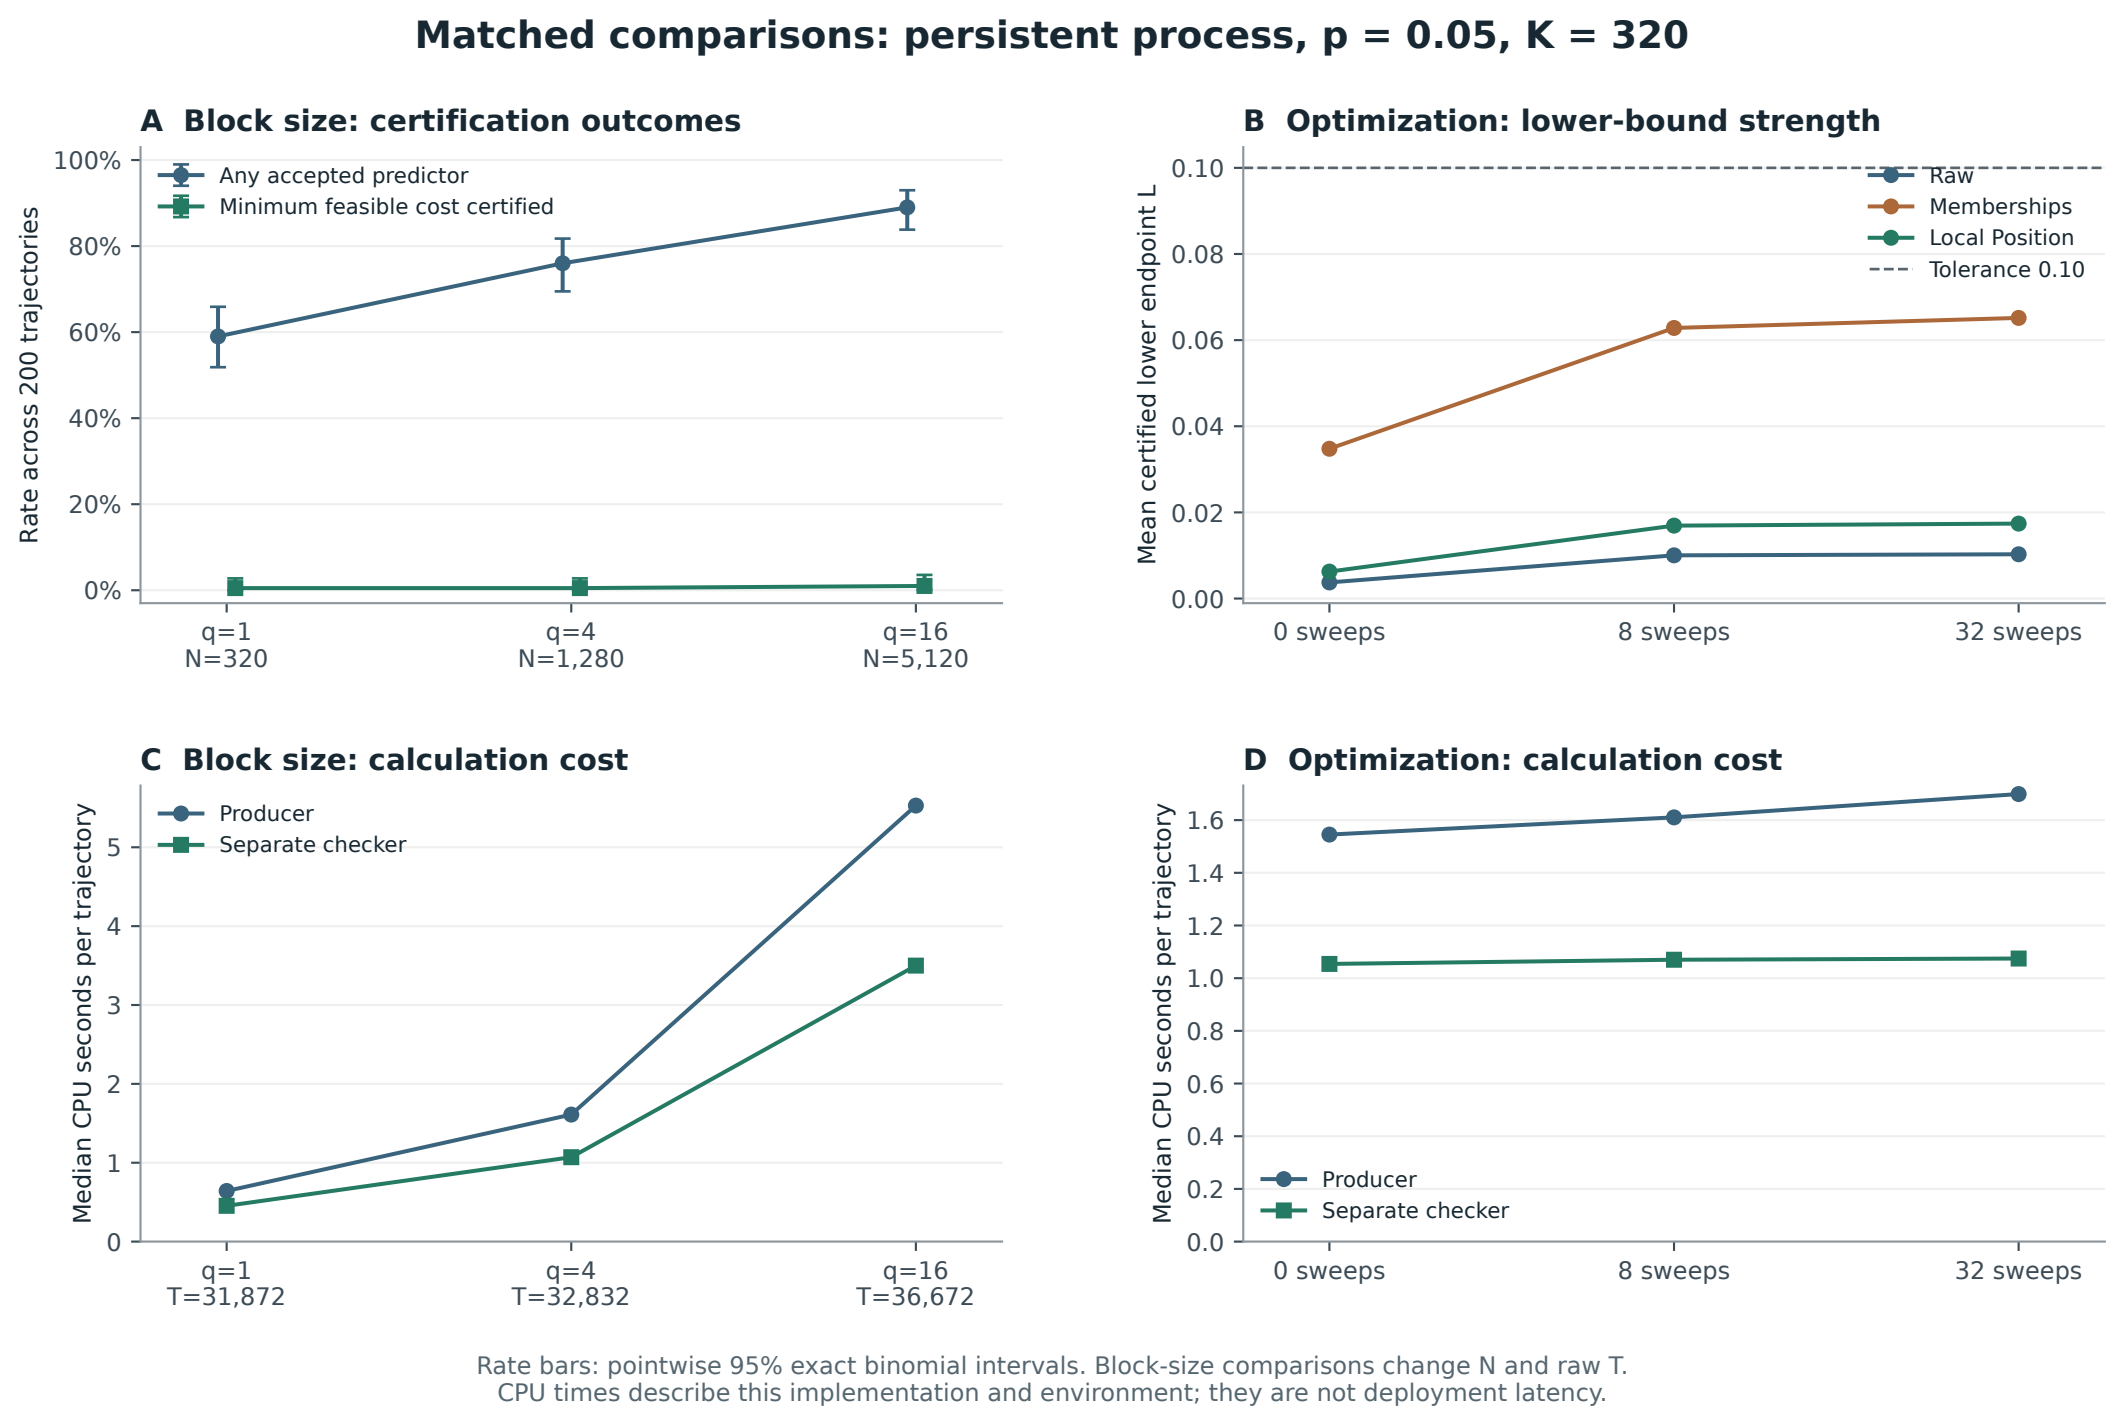
Matched comparisons: persistent process, p = 0.05, K = 320
A Block size: certification outcomes
Rate across 200 trajectories
100%
80%
60%
40%
20%
0%
Any accepted predictor
Minimum feasible cost certified
q=1
N=320
q=4
N=1,280
q=16
N=5,120
B Optimization: lower-bound strength
Mean certified lower endpoint L
0.10
0.08
0.06
0.04
0.02
0.00
Raw
Memberships
Local Position
Tolerance 0.10
0 sweeps
8 sweeps
32 sweeps
C Block size: calculation cost
Median CPU seconds per trajectory
0
1
2
3
4
5
Producer
Separate checker
q=1
T=31,872
q=4
T=32,832
q=16
T=36,672
D Optimization: calculation cost
Median CPU seconds per trajectory
0.0
0.2
0.4
0.6
0.8
1.0
1.2
1.4
1.6
Producer
Separate checker
0 sweeps
8 sweeps
32 sweeps
Rate bars: pointwise 95% exact binomial intervals. Block-size comparisons change N and raw T.
CPU times describe this implementation and environment; they are not deployment latency.